

Jupyter notebooks: How do kernels work internally?

Florian Thöle

PyData Zurich

September 29, 2017

About me

- PhD student in Materials Theory at ETH Zürich
- Python/Jupyter notebooks are my workhorse for everything around large-scale quantum-mechanical simulation programs
- Likes to teach – SWC workshops, workshops in my research area, ...

ETH zürich


software carpentry



@florian_thl



florian.thoele@gmail.com

Kernel Restarting



The kernel appears to have died. It will restart automatically.

OK

IPython vs. Jupyter

- Since version 4.0: Split between ipython and jupyter
- IPython: Provide a “more usable” python environment for interactive work
 - REPL and kernel mode
 - Environment on top of a python interpreter, such as CPython, PyPy, ..
- Jupyter: Client/server infrastructure, can use ipython as one possible kernel

IP[y]: IPython
Interactive Computing



Architecture of Jupyter



How does the server know about available kernels?

- KernelManager reads JSON files with kernelspec:

```
$ jupyter kernelspec list
```

```
Available kernels:
```

```
python3          /Users/thoelef/anaconda3/lib/python3.6/site-packages/ipykernel/resources
ir               /Users/thoelef/Library/Jupyter/kernels/ir
julia-0.4        /Users/thoelef/Library/Jupyter/kernels/julia-0.4
python2          /Users/thoelef/Library/Jupyter/kernels/python2
scala211         /Users/thoelef/Library/Jupyter/kernels/scala211
```

- nb_conda_kernels replaces the standard KernelManager with a custom one that “emulates” kernel confs for all conda environments

Example of KernelSpec file

- Usually, executes python program, but can be any other program that implements the messaging protocol

```
{
  "display_name": "Julia 0.4.5",
  "argv": [
    "/Applications/Julia-0.4.5.app/Contents/Resources/julia/bin/julia",
    "-i",
    "-F",
    "/Users/thoelef/.julia/v0.4/IJulia/src/kernel.jl",
    "{connection_file}"
  ],
  "language": "julia"
}
```

The connection file contains the sockets used for communication

```
{  
  "shell_port": 63454,  
  "iopub_port": 63455,  
  "stdin_port": 63456,  
  "control_port": 63457,  
  "hb_port": 63458,  
  "ip": "127.0.0.1",  
  "key": "c78c3588-0ab3-408c-8bc3-55b9cd9ee291",  
  "transport": "tcp",  
  "signature_scheme": "hmac-sha256",  
  "kernel_name": ""  
}
```


Jupyter message protocol

Messages are sent on five channels:

```
{  
  'header' : dict,  
  'msg_id' : str,  
  'msg_type' : str,  
  'parent_header' : dict,  
  'content' : dict,  
  'metadata' : dict,  
}
```

Shell

Control

IOPub

stdin

heartbeat

<http://jupyter-client.readthedocs.io/en/latest/messaging.html>

How to implement a custom kernel

- Write from scratch, have to start at socket level to send/receive messages
 - Can be done in any language
- Use ipython module that takes care of communication
 - Good if python as base layer works
 - Can be used to control REPLs

Live demo!

Examples of languages for which kernels are written

- IJulia, R, PySpark, JavaScript, Matlab, ...
- C, Fortran, C#, Scala, ...

<https://github.com/jupyter/jupyter/wiki/Jupyter-kernels>

https://github.com/dsblank/simple_kernel

<https://github.com/QuantStack/xeus>

Summary

- Versatile architecture of Jupyter allows custom kernels for a variety of languages
- Effort to support a new language low when it can be controlled from python (pexpect), higher when implemented from scratch in native language